

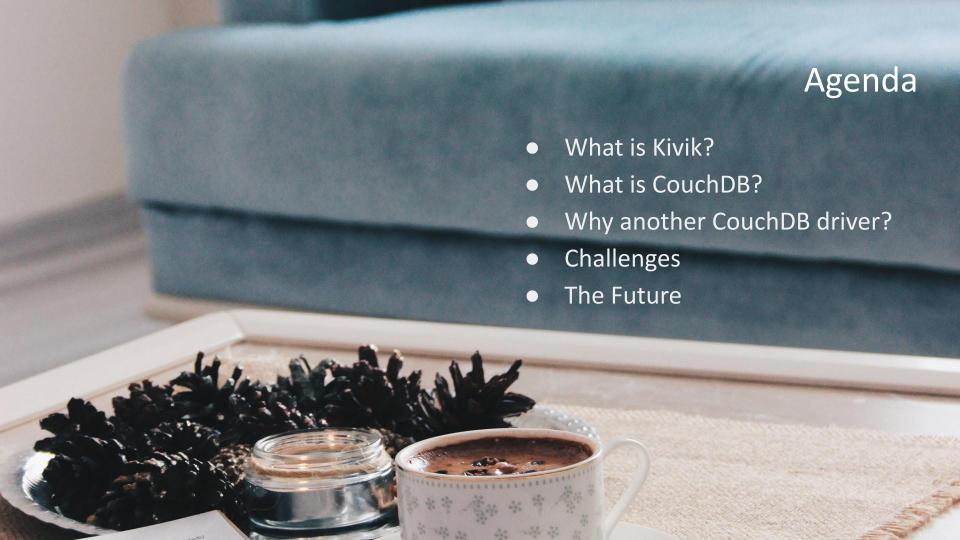


Hi, I'm Jonathan

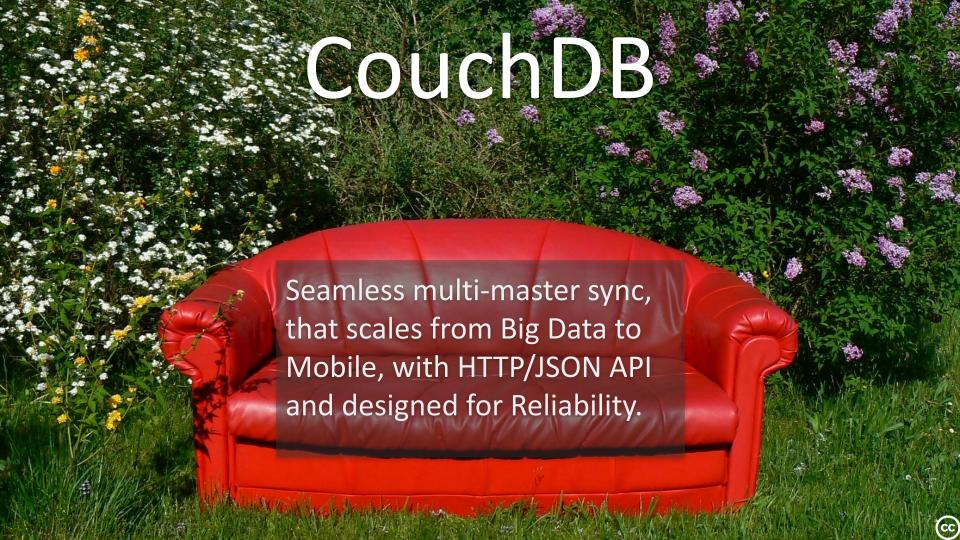
Go developer

Member CouchDB PMC

VP Engineering @ HUBUC









Consistency

Every read receives the most recent data, or an error

CP

Availability

Every read receives a response, but maybe outdated data

CA

Partition Tolerance

System operates even with failed nodes or network outages



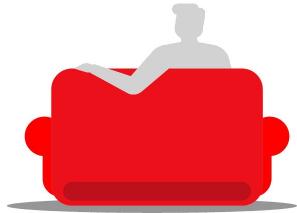


AP



The CouchDB API at a Glance

- CouchDB queries comprise HTTP verbs, specific endpoints, and an HTTP request
- CouchDB pretty strictly follows HTTP standards
 - Verbs: GET, HEAD, PUT, POST, DELETE (and COPY)
 - Request Headers: Accept, Content-Type, and a few others
 - Response Headers: Content-Type, ETag, and a few others
 - Status Codes: 200, 201, 202, 400, 404, etc
- Uses JSON
 - application/json for most requests and responses
 - Only rarely is anything else used







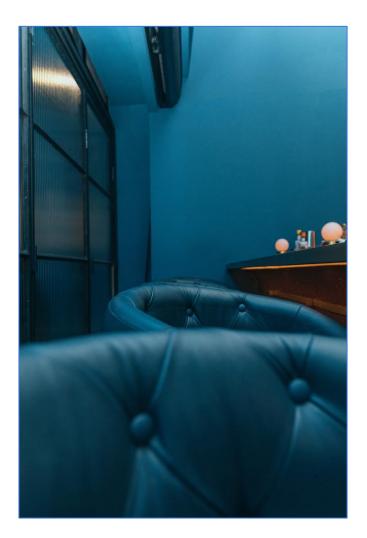
Design Goals

- JSON streaming support
- Open-source license
- Support for modern CouchDB features (i.e. find)
- A single API for CouchDB & PouchDB & more
- An idiomatic Go interface, balanced with idiomatic CouchDB
- Easy to proxy some CouchDB requests through a Go server
- An extensible & composable suite of tools
- A unique name that's easy to type and spell



My approach

- A "driver" model inspired by sql & sql/driver in stdlib
- Multiple back-ends:
 - CouchDB
 - PouchDB
 - Mock
 - Filesystem
 - In-memory
 - Proxy
- Multiple front-ends:
 - Your own software
 - A CLI
 - A server



But what to call it?

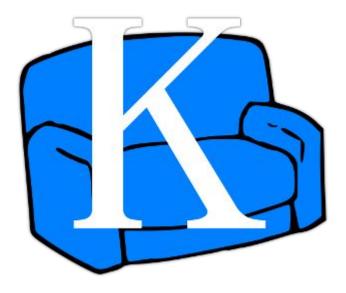
- Something related to couches
- Something modular
- Pick and choose the parts you want

This business model sounds familiar...



Kivik at a Glance

*kivik.Client for most server-scoped requests:



Kivik at a Glance

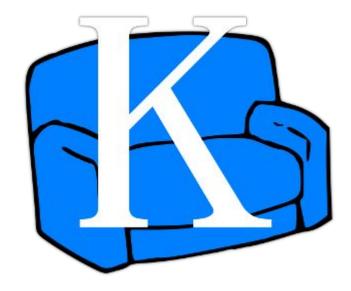
*kivik.DB for most database-scoped requests.

```
db := kivik.DB("my_db")

GET /my_db/_all_docs ⇒ db.AllDocs()

POST /my_db ⇒ db.CreateDoc()

PUT /my_db/doc1 ⇒ db.Put()
```





Sources of Errors

- Client errors (i.e. invalid URL)
- Driver errors (i.e. unsupported feature)
- Network errors (i.e. timeout)
- Server errors (i.e. document not found)

How can I support all of these in a consistent, and simple way?

Tried different approaches.



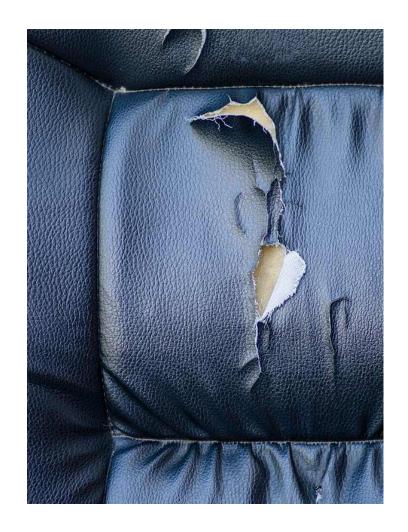
Current Approach

A single error type, and a helper function:

```
type Error struct {
    HTTPStatus int
    Message string
    Err error
}

func StatusCode(err error) int

And a few "fudged" status codes.
```



Checking an Error

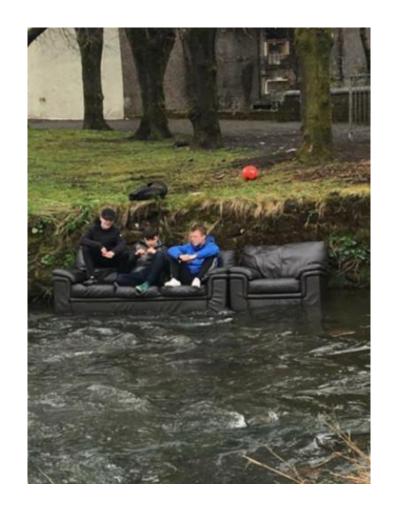
```
err := db.Get(context.TODO(), "docID").ScanDoc(&doc)
if kivik.StatusCode(err) == http.StatusNotFound {
    return
}
```

Or use errors.As/errors.Is to check for specific network failures, etc.



JSON Streaming

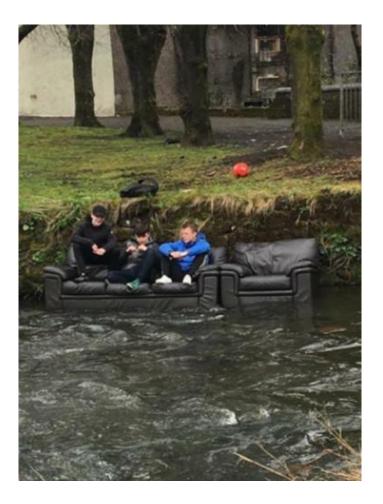
- No stdlib support
- json.Decoder works, but is very cumbersome



```
func (i *iter) next(row interface{}) error {
     if i.dec == nil {
            i.dec = json.NewDecoder(i.body)
           if err := i.begin(); err != nil {
                  return &kivik.Error{
                        HTTPStatus: 502, Err: err}
     err := i.nextRow(row)
     if err != nil {
           if err == io.EOF {
                  if e := i.finish(); e != nil {
                        err = e
                 return err
      return err
```



```
func (i *iter) begin() error {
     // consume the first '{'
     if e := consume(i.dec, json.Delim('{')); e != nil {
           return e
     for {
           key, err := nextKey(i.dec)
            if err != nil {
                  return err
            if key == i.expectedKey {
                 // Consume the first '['
                 return consume(i.dec, json.Delim('['))
            if err := i.parseMeta(key); err != nil {
                  return err
```



Re-Parsing JSON

Normally, the document revision is included in the ETag header of a response, making it easy to return without reading the entire response.

```
HTTP/1.1 200 OK
Cache-Control: must-revalidate
Content-Length: 660
Content-Type: application/json
Date: Tue, 13 Aug 2013 21:35:37 GMT
ETag: "12-151bb8678d45aaa949ec3698ef1c7e78"

doc := db.Get("my_doc")
fmt.Println(doc.Rev)
// 12-151bb8678d45aaa949ec3698ef1c7e78
```



Re-Parsing JSON

However, in some cases that's not possible, as when returning the document with extra data (which would make an ETag-based cache invalid).

In such a case, it is necessary to parse the response body to extract the revision.

```
HTTP/1.1 200 OK
Content-Length: 660
Content-Type: application/json
Date: Tue, 13 Aug 2013 21:35:37 GMT

{"_id":"my_doc","_rev":" 12-151bb8678d45aaa949ec3698ef1c
7e78",...
```



```
func extractRev(resp *http.Response) (string, error) {
      buf := &bytes.Buffer{}
      r := io.TeeReader(resp.Body, buf)
      defer func() {
            // Restore the original resp.Body
            resp.Body = struct {
                  io.Reader
                  io.Closer
            }{
                  Reader: io.MultiReader(buf, resp.Body),
                  Closer: resp.Body,
      }()
      rev, err := readRev(r)
      if err != nil {
            return "", err
      return rev, nil
```



```
func readRev(r io.Reader) (string, error) {
      dec := json.NewDecoder(r)
     tk, err := dec.Token()
      if err != nil {
           return "", err
      if tk != json.Delim('{') {
           return "", errors.New("wrong token")
      for dec.More() {
           tk, err = dec.Token()
            if err != nil {
                 return "", err
            if tk == "_rev" {
                 tk, err = dec.Token()
                  if err != nil {
                       return "", err
                  if value, ok := tk.(string); ok {
                       return value, nil
                  return "", errors.New("invalid JSON")
      return "", errors.New("_rev key not found")
```



History

June 2017: 1.0 release w/ CouchDB & PouchDB drivers

Also partial/experimental support for:

- Filesystem, in-memory, and proxy drivers
- Server
- Feb 2020: 2.0/30 release w/ kivikmock for testing
- Dec 2020: CLI tool
 - Extensive filesystem support (without indexes)
 - Allows for easy disk-to-Couch or Couch-to-disk replication (analogous to pg_restore and pg_dump)



Future

- ???: 4.0 release (~75% complete)
 - Improved error handling
 - API improvements, particularly for multi-document responses
- Add support for full-text search, added in CouchDB 3.0
- Support incremental replication in the CLI tool (useful for incremental backup of a CouchDB server)
- Build Go indexing support, to allow indexing in the FS and in-memory backends
- Flesh out the server tool, possibly merge it with the CLI tool
- Rename package import path to kivik.io?
- Re-consolidate drivers into main package?





Other Resources

CouchDB website: https://couchdb.apache.org/

Kivik website: http://kivik.io/

YouTube Video: Advanced JSON Handling in Go

Very incomplete eBook: <u>Data Serialization in Go</u>

New YouTube Channel: **Boldly Go**



Slides at:

https://jhall.io/kivik-story