

HTTPq: DDoS-as-a-Service

How MessageBird delivers millions of webhooks daily



Disclaimer

MessageBird **condemns** any malicious attack (DDoS or not) and works together both with customers and engineers everywhere to make sure security and operations are taken **seriously**. Any situations presented here are anecdotal scenarios gathered across multiple years operating a system designed to perform as fast as it can for customers who need as many webhooks as they'd need.



About me


Brazilian bird 🇧🇷

(slightly obsessed) cat person 🐱

On the (tech) road for 17 years 🦉

Tech geek since forever 🎮

Business Infra Team Lead @ MessageBird 🐦




Business Infra

Manages product infrastructure services

Authentication, queues, preferences

Flamingos and cats everywhere



About MessageBird

Communications Platform as a Service (CPaaS)

Global scale with big customers

A lot of challenges scaling and innovating

Fun place to work (we're hiring!)



Let's talk about HTTPq



```

// RunWorker is responsible for fetching a job from beanstalk and passing that
// job to handleJob().
func RunWorker() {
    for {
        handleJob(<-consumerPool.C)
    }
}

// handleJob extracts the Request ID from the job's JSON data, retrieves the
// Request from the database and tries to perform an HTTP(S) request.
func handleJob(job *beanstalk.Job) {
    var (
        request    *httpqueue.Request
        requestJob beanstalkRequest
    )

    // Extract the ID from the job's JSON data and fetch the Request object.
    err := json.Unmarshal(job.Body, &requestJob)
    if err == nil {
        request, err = httpqueue.RequestWithID(requestJob.ID)
    }
    if err != nil {
        log.Printf("Unable to fetch request for beanstalk job %d: %s", job.ID, err)
        if err := job.Bury(); err != nil {
            log.Printf("Unable to bury beanstalk job %d: %s", job.ID, err)
        }
    }

    return
}

Debug.Printf("%d: Worker parsing request", request.ID)

// If the request is ready, perform the actual HTTP request.
if request.IsReady() {
    Debug.Printf("%d: Worker performing request", request.ID)

    if err := request.Do(); err != nil && request.Status == httpqueue.StatusFailed {
        log.Printf("Request %d failed after %d attempts: %s", request.ID, request.Attempts, err)
    }

    if err := request.Save(); err != nil {
        log.Printf("Unable to save request %d: %s", request.ID, err)
    }
}

// If the current state requires a callback, do it here.
if request.ShouldCallback() {
    Debug.Printf("%d: Worker performing callback for status '%s'", request.ID, request.Status)

    if err := request.DoCallback(producerPool); err != nil {
        if request.IsFinished() {
            log.Printf("Request %d was unable to perform final callback: %s", request.ID, err)
            if err := job.Bury(); err != nil {
                log.Printf("Unable to bury beanstalk job %d: %s", job.ID, err)
            }
        }
        return
    }
}

if err := request.Save(); err != nil {
    log.Printf("Unable to save request %d: %s", request.ID, err)
}

// If there is no more work for this request, delete the job.
if request.IsFinished() && !request.ShouldCallback() {
    Debug.Printf("%d: Worker deleting beanstalk job", request.ID)
}

```

What is HTTPq?

Developed in 2016 by our dinosaurs

Queue for webhooks written in Go

Connects MessageBird services & customers

Deliver consistently across multiple clients

MessageBird delivers a lot of webhooks daily



Question: how many webhooks a day on average?

- a) Tens of thousands
- b) Hundreds of thousands
- c) Tens of millions
- d) Hundreds of millions




Question: how many webhooks a day on average?

- a) Tens of thousands
- b) Hundreds of thousands
- c) Tens of millions
- d) **Hundreds of millions** ✓



Webhooks in August 2022



Breaking down HTTPq



Breaking down HTTPq

Self-contained single binary

KISS principle

Secure by default

Three main components in the service





100

Continue

Breaking down HTTPq: API

HTTP and GRPC APIs available

Single endpoint which users POST requests

Runs on Go's standard HTTP mux

Globally available (thanks to Google GLB)

Horizontal Pod Autoscaling to adapt to traffic

99.95% submissions under 250ms internal SLO



Breaking down HTTPq: Beanstalk

Open source queue manager

Simple protocol

Delaying and priority commands

Fast, single-threaded process

Nice open source Go client written by one of our birds

<https://github.com/beanstalkd/beanstalkd>

<https://github.com/prep/beanstalk>





Breaking down HTTPq: Worker

Configurable number of workers

Handles the heavy lifting for webhooks

Main source of our issues in the past

Socket ports hungry

HPA to adapt to traffic

99.95% deliveries under 500ms internal
SLO

(measured by the moment the webhook request comes in)



Lifecycle of a webhook

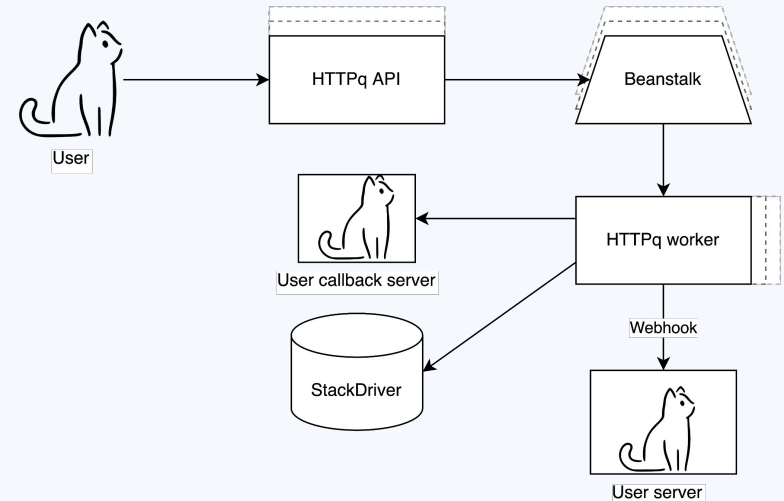
Schedule webhook deliveries using HTTPq's API

Webhooks are queued on beanstalk

Workers pick up pending webhook requests from beanstalk

Request arrives to target server

(Optional) Callback request is queued



Operating & maintaining HTTPq



Scaling HTTPq



Horizontal pod autoscaling

Multi-region installation

Global load balancing

Dedicated nodes

Finding the resource quotas sweet spot



HTTP2 client issues

Weird behaviour since go1.9 – up until go1.17

Caused process hugging

Lots of sleepless nights

<https://github.com/golang/go/issues/32388>

Small patch to fix that



No transactional databases! Why?

Scaling is painful

Multi-region database can be expensive...

... and might not work as expected

In the end, transactional databases are not really needed



Embrace audit logging



Structured logging

Logs are non-blocking

Cloud Logging is your friend...


... but beware of costs!



Question: how much it costs per month?



- a) Hundreds of thousands euros
- b) Tens of thousands euros
- c) A few thousands euros
- d) Less than a thousand euros




Question: how much it costs per month?

- a) Hundreds of thousands euros
- b) Tens of thousands euros
- c) **A few thousands euros** ✓
- d) Less than a thousand euros



Costs in the last 60 days



Conclusion/main takeaways

Design a system to solve your current problems

Invest time and money scaling when needed

Services doesn't need to be expensive to do a lot

Issues exists everywhere

(even in Golang's codebase)

Be careful what you wish for!



Q&A

Ask us anything



Thanks for your attention!

