

A small horizontal bar with a teal segment on the left and an orange segment on the right.

Concurrency patterns

Ilija Matoski
ilijamt@gmail.com

2022-09-28



<https://matoski.com>




<https://linkedin.com/ilijamt>



<https://github.com/ilijamt>

-
- Concurrency is the composition of independently executing things (functions or processes in the abstract).
 - Parallelism is the simultaneous execution of multiple things (possibly related, possibly not)
 - Concurrency is dealing with a lot of things at once.
 - Parallelism is doing a lot of things at once.






Building blocks of Concurrency in Go (goroutines)

- Channels
- Buffered Channels
- Select statement
- Wait groups
- Mutex
- [sync - Package](#)
- [runtime - Go Package](#)



goroutines

- Fork/Join model
- Not necessarily running in parallel
- Have the same address space
- goroutines as light-weight threads (few kb of memory)
- Have their own call stack which can grow and shrink dynamically.





channels

- Based on Communicating Sequential Processes put down by Hoare in 1978
- Golang makes use of channels to achieve communication between concurrent processes
- Don't communicate by sharing memory; share memory by communicating





channels

`close(channelName)`

`channelName := make(chan datatype)`

`channelName := make(chan datatype, capacity)`

`channelName <- data` <- Sending on a channel

`data := <- channelName` <- Receiving on a channel





channels - Directions

```
func ping(pings chan<- string, msg string) {  
    pings <- msg  
}
```

```
func pong(pings <-chan string, pongs chan<- string) {  
    msg := <-pings  
    pongs <- msg  
}
```

```
pings := make(chan string, 1)  
pongs := make(chan string, 1)  
ping(pings, "passed message")  
pong(pings, pongs)  
fmt.Println(<-pongs)
```





select

```
select {  
  case channel operation:  
    statement(s);  
    . //more cases  
    .  
    .  
  default : //Optional  
    statement(s);  
}
```





select

The select statement blocks the code and waits for multiple channel operations simultaneously.

- Default Case
- Nil Channel
- Empty Select

```
select {  
    case channel operation:  
        statement(s);  
        . //more cases  
        .  
        .  
    default : //Optional  
        statement(s);  
}
```





select - Default Case

- Executes when every other send/receive operation is blocked.
- If we have more than one send/receive operation ready at the same time? The select statement chooses one at random

```
select {  
  case msg1 := <-channel1:  
    fmt.Println(msg1)  
  case msg2 := <-channel2:  
    fmt.Println(msg2)  
  default:  
    fmt.Println("No channel is ready")  
}
```





select - Nil channel

- Nil channel will block the select statement forever, giving an error stating a deadlock
- To disable the send/receive operations of a channel in select statement, we can use nil channels
- To avoid deadlock we need to add a default case or a case that is not blocked

```
ch1 := make(chan string)
ch1 = nil
ch2 := make(chan string)
```

```
select {
case msg1 := <-ch1:
    fmt.Println(msg1)
case msg2 := <-ch2:
    fmt.Println(msg2)
default:
    fmt.Println("No channel is ready")
}
```





select - Empty Select

- Blocks execution of code on the current goroutine
- **for { }** will cause the CPU% to max, and the process's STATE will be running
- **select { }**, will not cause the CPU% to max, and the process's STATE will be sleeping

```
package main
```

```
func main() {  
    select {}  
}
```

```
$ go run empty-select.go
```

```
fatal error: all goroutines are asleep - deadlock!
```

```
goroutine 1 [select (no cases)]:
```

```
main.main()
```

```
  /tmp/run.go:4 +0x17
```

```
exit status 2
```





sync

- Cond
- Mutex
- Once
- Pool
- RWMutex
- WaitGroup
- Semaphore





sync.Cond

- Implements a condition variable which indicates the goroutines which are waiting for an event or want to announce an event
- The thing to note is that if `cond.Wait()` before `cond.Broadcast()` then the listeners will keep on waiting forever.





sync.Cond

```
var wg = new(sync.WaitGroup)
var vals = make(map[string]int)
wg.Add(5)
```

```
m := sync.Mutex{}
cond := sync.NewCond(&m)
```

```
go listen("lis1", vals, cond, wg)
go listen("lis2", vals, cond, wg)
```

```
go broadcast("b1", vals, cond, wg)
```

```
wg.Wait()
```

```
func listen(name string, a map[string]int, c *sync.Cond, wg *sync.WaitGroup) {
    defer wg.Done()
    c.L.Lock()
    c.Wait()
    fmt.Println(name, " val:", a["T"])
    c.L.Unlock()
}
```

```
func broadcast(name string, a map[string]int, c *sync.Cond, wg *sync.WaitGroup) {
    defer wg.Done()
    time.Sleep(time.Second)
    c.L.Lock()
    a["T"] = 25
    c.Broadcast()
    c.L.Unlock()
}
```






sync.Mutex

- A mutex, or a mutual exclusion, prevents other processes from entering a critical section of data while a process occupies it.

```
mu := sync.Mutex{}
mu.Lock()
go func() {
    mu.Lock() // <- This blocks until mu.Unlock() is called
    fmt.Println("I am in the goroutine")
    mu.Unlock()
}()
fmt.Println("I am in main routine")
mu.Unlock()
time.Sleep(time.Second)
```

sync.RWMutex

- Stands for Reader/Writer mutual exclusion and is essentially the same as Mutex, but it gives the lock to more than one reading process or just a writing process





sync.Once

- Is an object that performs an action only once



```
var once sync.Once
onceBody := func() {
    fmt.Println("Only once")
}
done := make(chan bool)
for i := 0; i < 10; i++ {
    go func() {
        once.Do(onceBody)
        done <- true
    }()
}
for i := 0; i < 10; i++ {
    <-done
}
```



sync.Pool

- Is a collection of temporary objects which can be accessed and saved by many goroutines simultaneously
- Any item stored in the Pool may be removed automatically at any time without notification. If the Pool holds the only reference when this happens, the item might be deallocated.
- A Pool is safe for use by multiple goroutines simultaneously.
- Pool's purpose is to cache allocated but unused items for later reuse, relieving pressure on the garbage collector.
- An example of good use of a Pool is in the `fmt` package, which maintains a dynamically-sized store of temporary output buffers





sync.WaitGroup

- Blocks a program and waits for a set of goroutines to finish before moving to the next steps of execution



```
var wg sync.WaitGroup
```

```
wg.Add(2)
```

```
go func() {  
    defer wg.Done()  
    // Do work.  
    r := time.Duration(rand.Intn(100)) * time.Microsecond  
    time.Sleep(r)  
    fmt.Printf("Worked for %s\n", r)
```

```
}()
```

```
go func() {  
    defer wg.Done()  
    // Do work.  
    r := time.Duration(rand.Intn(100)) * time.Microsecond  
    time.Sleep(r)  
    fmt.Printf("Worked for %s\n", r)
```

```
}()
```

```
wg.Wait()  
fmt.Println("Finished")
```





sync.Semaphore

- Use them as a Worker Pool
- Binary Semaphore
- Counting Semaphore





runtime.GOMAXPROCS(n int) int

- GOMAXPROCS allows us to set the maximum number of CPUs that can be executed simultaneously.
- Calling the function will set the number of CPUs to be **n** but will return the previous value set for the number of the CPUs.
- Calling the function with **n < 1** then we return the previous setting

```
fmt.Printf("GOMAXPROCS is %d\n", runtime.GOMAXPROCS(3)) // previous setting  
fmt.Printf("GOMAXPROCS is %d\n", runtime.GOMAXPROCS(0)) // will return 3
```





`runtime.NumCPU() int`

- Returns the number of logical CPUs that can be used by the current process.





Benchmarks - synchronization of values

- No synchronization
- Using atomics
- Mutex
- Buffered channels





Benchmarks - No synchronization

```
var j int64
```

```
func Benchmark_Int64(b *testing.B) {  
    for i := 0; i < b.N; i++ {  
        j++  
    }  
}
```

| | i7-8550U | 5950X |
|-------|-------------|-----------|
| ops | 870 million | 1 billion |
| ns/op | 1.409 | 0.2247 |





Benchmarks - Using atomics

```
var j int64
```

```
func Benchmark_Int64(b *testing.B) {  
    for i := 0; i < b.N; i++ {  
        atomic.AddInt64(&j, 1)  
    }  
}
```

| | i7-8550U | 5950X |
|--------------|-----------------|--------------|
| ops | 183 million | 737 million |
| ns/op | 6.638 | 1.638 |





Benchmarks - Using Mutex

```
var j int64
var m sync.Mutex

func Benchmark_Int64(b *testing.B) {
    for i := 0; i < b.N; i++ {
        m.Lock()
        j++
        m.Unlock()
    }
}
```

| | i7-8550U | 5950X |
|-------|------------|-------------|
| ops | 92 million | 306 million |
| ns/op | 13.07 | 3.888 |





Benchmarks - Channels

```
var j int64
var c make(chan interface{}, 1)

func Benchmark_Int64(b *testing.B) {
    for i := 0; i < b.N; i++ {
        c <- nil
        j++
        <- c
    }
}
```

| | i7-8550U | 5950X |
|-------|------------|------------|
| ops | 30 million | 57 million |
| ns/op | 42.70 | 20.18 |





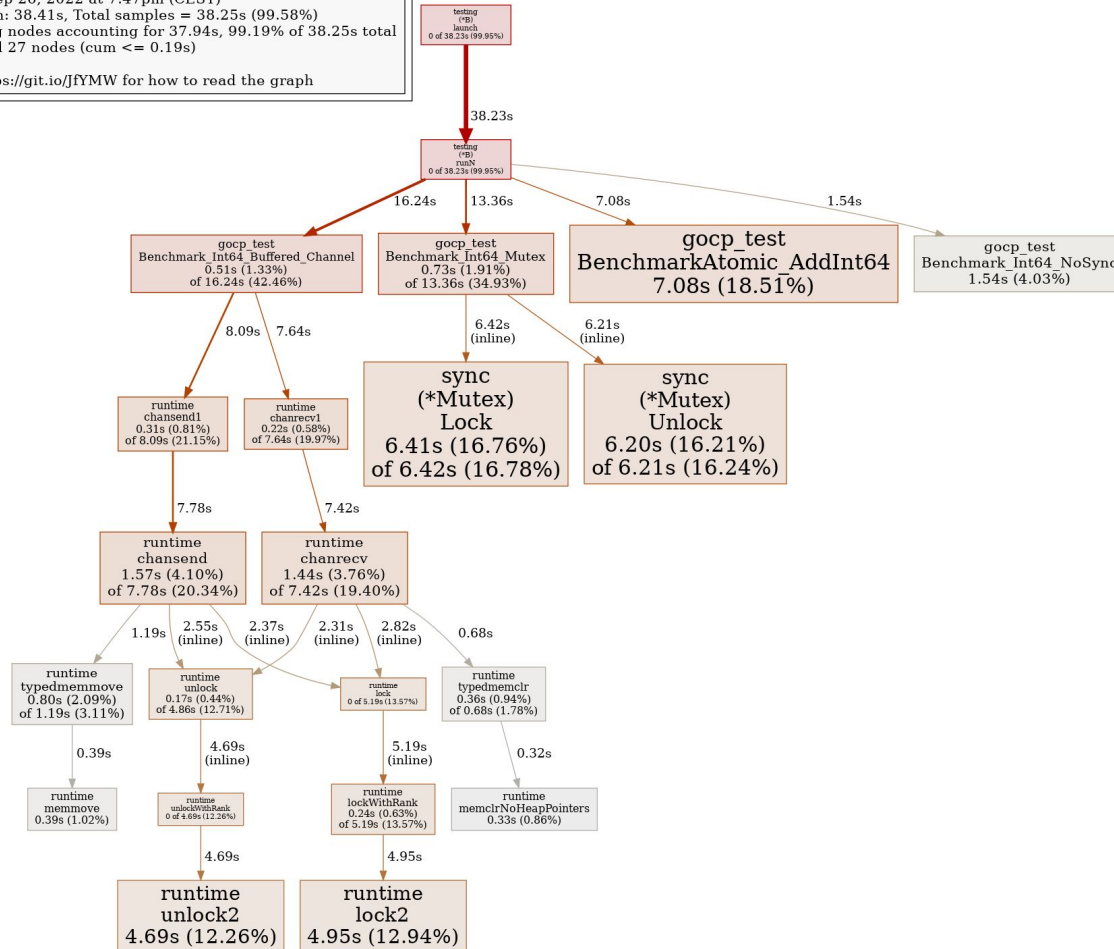
Benchmarks

| | No synchronization | | atomics | | Mutex | | Channels | |
|-------------|-----------------------|--------|----------|-------|----------|-------|----------|-------|
| | i7-8550U | 5950X | i7-8550U | 5950X | i7-8550U | 5950X | i7-8550U | 5950X |
| ops/million | 870 | 1000 | 183 | 737 | 92 | 306 | 30 | 57 |
| ns/op | 1.409 | 0.2247 | 6.638 | 1.638 | 13.07 | 3.888 | 42.7 | 20.18 |




File: gocp.test
Type: cpu
Time: Sep 26, 2022 at 7:47pm (CEST)
Duration: 38.41s, Total samples = 38.25s (99.58%)
Showing nodes accounting for 37.94s, 99.19% of 38.25s total
Dropped 27 nodes (cum <= 0.19s)

See <https://git.io/JfYMW> for how to read the graph



Patterns in Go

- Generator
- Fan-In, Fan-Out
- Sequencing
- Range and Close
- For-Select
- Quit Channel
- Timeout using Select
- Channel Ring Buffer
- Barriers





Generator

- Generators return the next value in a sequence each time they are called.
- Each value is available as an output before the generator computes the next value.

```
func fibonacci(n int) chan int {  
    ch := make(chan int)  
    go func() {  
        k := 0  
        for i, j := 0, 1; k < n; k++ {  
            ch <- i  
            i, j = i+j, i  
        }  
        close(ch)  
    }()  
    return ch  
}
```





For-Select

- Select statement lets a goroutine wait on multiple operations
- Select blocks until one of the cases can run
- A random one is chosen if multiple are ready
- If there is a default, it will execute on every loop without blocking the loop and waiting on a value from a channel

```
for {  
    select {  
        case i := <- c1:  
            // use i  
        case a := <- c2  
            // use a  
        default:  
            // executes on every turn  
        }  
    }
```



Range and Close

- Only the sender should close a channel, never the receiver
- Sending on a closed channel will cause a panic
- Closing is only necessary when the receiver must be told there are no more values coming, such as to terminate a range loop.



```
func fibonacci(n int, c chan int) {  
    x, y := 0, 1  
    for i := 0; i < n; i++ {  
        c <- x  
        x, y = y, x+y  
    }  
    close(c)  
}
```

```
c := make(chan int, 10)  
go fibonacci(cap(c), c)  
for i := range c {  
    fmt.Println(i)  
}
```





Timeout using Select

```
for {
    select {
        case s := <-ch:
            fmt.Println(s)
            return
        case <-time.After(time.Second):
            fmt.Println("bail!")
            return
    }
}
```





Quit Channel

```
channel := make(chan string)
quit := make(chan string)

for i := 0; i < 6; i++ {
    go Race(channel, quit, i)
}

for {
    select {
    case raceUpdates := <-channel:
        fmt.Println(raceUpdates)
    case winnerAnnouncement := <-quit:
        fmt.Println(winnerAnnouncement)
        return
    }
}
```

```
func Race(channel, quit chan string, i int) {
    channel <- fmt.Sprintf("Mouse %d started!", i)
    for {
        rand.Seed(time.Now().UnixNano())
        time.Sleep(time.Duration(rand.Intn(500)+500) * time.Millisecond)
        quit <- fmt.Sprintf("Mouse %d reached the finishing line!", i)
    }
}
```

```
Mouse 5 started!
Mouse 1 started!
Mouse 0 started!
Mouse 3 started!
Mouse 2 started!
Mouse 4 started!
Mouse 5 reached the finishing line!
```






Channel Ring Buffer

- Connect two buffered channels that forward messages from the incoming channel to the outgoing channel
- If a message cannot be placed on the outgoing channel, drop the oldest message in the outgoing channel
- Slow consumers might lose messages (their oldest messages), but it will never block





Channel Ring Buffer

```
type RingBuffer struct {
    inputCh <-chan int
    outputCh chan int
}

func (r *RingBuffer) Run() {
    for v := range r.inputCh {
        select {
        case r.outputCh <- v:
        default:
            <-r.outputCh
            r.outputCh <- v
        }
    }
    close(r.outputCh)
}

func main() {
    in := make(chan int)
    out := make(chan int, 5)
    rb := &RingBuffer{in, out}
    go rb.Run()

    for i := 0; i < 10; i++ {
        in <- i
    }

    close(in)

    for res := range out {
        fmt.Println(res)
    }
}
```

```
go run ring-buffer.go
```

```
4
5
6
7
8
9
```




Fan-In, Fan-Out

- It's a way to converge and diverge data into a single data stream from multiple streams or from one stream to multiple streams or pipelines
- Fan-In refers to a technique in which you join data from multiple inputs into a single entity.
- Fan-Out means to divide the data from a single source into multiple smaller chunks. In this lesson, we'll learn how to make use of both these techniques.



Fan-in



Fan-out





Questions

