

BoltDB

An embedded K/V database for Go

Huseyin Ekemen

Golang & DevOps Developer

<https://www.nextpax.com>



Co-Founder @ pdaccess:

<https://www.pdaccess.com>



5 years in Golang

19 Years with Java

What we have been looking for:

- Single File Database (like sqlite)
- Should be pure go based.
- Has to consist about crash and concurrency cases.
- Reasonably good performance.(Both write & read)

BoltDB:

- LMDB(The Lightning Memory-Mapped Database)

- It is very small code base.
- It is embedded. Run inside your binary
- Single File DataStore
- Simple and stable api.
- ACID based (lock-free MVCC

using a single writer and multiple
readers)

MMAP(2) Linux Programmer's Manual MMAP(2)

NAME [top](#)

mmap, munmap – map or unmap files or devices into memory

SYNOPSIS [top](#)

```
#include <sys/mman.h>

void *mmap(void *addr, size_t length, int prot, int flags,
           int fd, off_t offset);
int munmap(void *addr, size_t length);
```

See NOTES for information on feature test macro requirements.

DESCRIPTION [top](#)

mmap() creates a new mapping in the virtual address space of the calling process. The starting address for the new mapping is specified in *addr*. The *length* argument specifies the length of the mapping (which must be greater than 0).

Concept:

- File == Database
 - Bucket == Table
 - Key-Value == map[[]byte][]byte
 - Transaction == Lock
-
- ReadOnly: Multiple read-only transaction is ok
 - ReadWrite: Only one readwrite transaction is allowed.
 - Batch: Multiple write together.

DB: Open

```
db, err := bolt.Open("my.db", 0600,  
&bolt.Options{Timeout: 1 * time.Second})  
  
// my.db is the filename of the database.  
// there are different options for bolt.Open  
// After usage you should close the db instance  
// db.Close()
```

Transaction: Write

```
db.Update(func(tx *bolt.Tx) error {
    b, err := tx.CreateBucket([]byte("MyLogBucket"))
    if err != nil {
        return fmt.Errorf("create bucket: %s", err)
    }
    return nil
})

// There is a method named CreateBucketIfNotExists()
```

Transaction: Write

```
db.Update(func(tx *bolt.Tx) error {  
    b := tx.Bucket([]byte("MyLogBucket"))  
    err := b.Put([]byte("user"), []byte("user1"))  
    return err  
})
```

```
// For Auto increment
```

```
// id, _ := b.NextSequence()
```


Transaction: Read

```
db.View(func(tx *bolt.Tx) error {  
    b := tx.Bucket([]byte("MyLogBucket"))  
    v := b.Get([]byte("user"))  
    fmt.Printf("The User is: %s\n", v)  
    return nil  
})
```

Transaction: Read

```
db.View(func(tx *bolt.Tx) error {
    // Assume bucket exists and has keys
    b := tx.Bucket([]byte("MyLogBucket"))

    c := b.Cursor()

    for k, v := c.First(); k != nil; k, v = c.Next() {
        fmt.Printf("key=%s, value=%s\n", k, v)
    }

    return nil
})

// The following functions are available on the cursor:
// First()  Move to the first key.
// Last()   Move to the last key.
// Seek()   Move to a specific key.
// Next()   Move to the next key.
// Prev()   Move to the previous key.
```

Transaction: Batch

```
db.Batch(func(tx *bolt.Tx) error {  
    b := tx.Bucket([]byte("MyLogBucket"))  
    b.Put([]byte("user"), []byte("user1"))  
    b.Put([]byte("device"), []byte("device1"))  
  
    return err  
})  
  
// Each update require file write  
// Single batch can help here.
```

More

- Dump stats easily
- Data is stored as []byte type.
- You can open database as read-only mode many times.
- Only one write-based access to the database file.
- Single func backup: `_, err := tx.WriteTo(w)`
- You can use `Seek()` for Range & prefix scan
- `ForEach()` to loop over a bucket.

```
// Grab the initial stats.  
prev := db.Stats()  
  
for {  
    // Wait for 10s.  
    time.Sleep(10 * time.Second)  
  
    // Grab the current stats and diff them.  
    stats := db.Stats()  
    diff := stats.Sub(&prev)  
  
    // Encode stats to JSON and print to STDERR.  
    json.NewEncoder(os.Stderr).Encode(diff)  
  
    // Save stats for the next loop.  
    prev = stats  
}
```

Links

- <https://github.com/boltdb/bolt> main project. It is abundant but still ok.
- <https://github.com/etcd-io/bbolt> successor of boltdb with same interface. You can continue with this one. (has live project.)
- <http://www.lmdb.tech/media/20130406-LOADays-LMDB.pdf> LMDB concept.